

Institute of Information Technology (IIT)
Jahangirnagar University



Lab Report: 04

Submitted by:

Name: Zannat Hossain Tamim

Roll No:1970

Lab Date:08.08.23

Submission Date: 10.08.23

Lab Report # Day 04

Query 1:

Read CSV file and print the top 5 rows.

Clause:

```
import numpy as np
import pandas as pd
df=pd.read_csv('housing.csv')
df.head()
```

Result :

Out[2]:

	0.00632	18.00	2.310	0	0.5380	6.5750	65.20	4.0900	1	296.0	15.30	396.90	4.98	24.00
0	0.02731	0.00	7.070	0	0.4690	6.4210	78...							
1	0.02729	0.00	7.070	0	0.4690	7.1850	61...							
2	0.03237	0.00	2.180	0	0.4580	6.9980	45...							
3	0.06905	0.00	2.180	0	0.4580	7.1470	54...							
4	0.02985	0.00	2.180	0	0.4580	6.4300	58...							

Query 2:

To create a list of column names and assign them to the variable column and print top 5 rows.

Clause:

```
column= ['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX', 'PTRATIO',  
'B', 'LSTAT', 'MEDV']  
data = pd.read_csv('housing.csv', header=None, delimiter=r"\s+", names=column)  
data.head()
```

Result :

```
In [6]: data.head()
```

Out[6]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	5.33	36.2

Query 3:

To get the shape of the DataFrame

Clause:

```
data.shape
```

Result :

```
In [7]: data.shape
```

Out[7]: (506, 14)

Query 4:

To get the information of the DataFrame

Clause:

```
data.info()
```

Result :

```
In [8]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   CRIM        506 non-null   float64
 1   ZN          506 non-null   float64
 2   INDUS       506 non-null   float64
 3   CHAS        506 non-null   int64  
 4   NOX         506 non-null   float64
 5   RM          506 non-null   float64
 6   AGE         506 non-null   float64
 7   DIS         506 non-null   float64
 8   RAD         506 non-null   int64  
 9   TAX         506 non-null   float64
10  PTRATIO     506 non-null   float64
11  B           506 non-null   float64
12  LSTAT       506 non-null   float64
13  MEDV        506 non-null   float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

Query 5:

To calculate and print a summary of statistical data for each column in the DataFrame

Clause:

```
data.describe()
```

Result :

```
In [9]: data.describe()
```

```
Out[9]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	L
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.237154	18.455534	356.674032	12.658292
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116	2.164946	91.294864	7.146154
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000	0.320000	1.730000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000	17.400000	375.377500	6.969900
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.000000	19.050000	391.440000	11.362300
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000	20.200000	396.225000	16.965000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000	396.900000	37.970000

Query 6:

To count the number of missing values in each column of a DataFrame

Clause:

```
data.isnull().sum()
```

Result:

```
In [10]: data.isnull().sum()
```

```
Out[10]: CRIM      0
          ZN        0
          INDUS    0
          CHAS     0
          NOX      0
          RM       0
          AGE      0
          DIS      0
          RAD      0
          TAX      0
          PTRATIO  0
          B        0
          LSTAT    0
          MEDV     0
          dtvpe: int64
```

Query 7:

Calculate the correlation coefficient between each pair of columns in the DataFrame and store it in the variable cor

Clause:

```
cor=data.corr()
print(cor)
```

Result:

```
print(cor)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	\
CRIM	1.000000	-0.200469	0.406583	-0.055892	0.420972	-0.219247	0.352734	
ZN	-0.200469	1.000000	-0.533828	-0.042697	-0.516604	0.311991	-0.569537	
INDUS	0.406583	-0.533828	1.000000	0.062938	0.763651	-0.391676	0.644779	
CHAS	-0.055892	-0.042697	0.062938	1.000000	0.091203	0.091251	0.086518	
NOX	0.420972	-0.516604	0.763651	0.091203	1.000000	-0.302188	0.731470	
RM	-0.219247	0.311991	-0.391676	0.091251	-0.302188	1.000000	-0.240265	
AGE	0.352734	-0.569537	0.644779	0.086518	0.731470	-0.240265	1.000000	
DIS	-0.379670	0.664408	-0.708027	-0.099176	-0.769230	0.205246	-0.747881	
RAD	0.625505	-0.311948	0.595129	-0.007368	0.611441	-0.209847	0.456022	
TAX	0.582764	-0.314563	0.720760	-0.035587	0.668023	-0.292048	0.506456	
PTRATIO	0.289946	-0.391679	0.383248	-0.121515	0.188933	-0.355501	0.261515	
B	-0.385064	0.175520	-0.356977	0.048788	-0.380051	0.128069	-0.273534	
LSTAT	0.455621	-0.412995	0.603800	-0.053929	0.590879	-0.613808	0.602339	
MEDV	-0.388305	0.360445	-0.483725	0.175260	-0.427321	0.695360	-0.376955	


	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
CRIM	-0.379670	0.625505	0.582764	0.289946	-0.385064	0.455621	-0.388305
ZN	0.664408	-0.311948	-0.314563	-0.391679	0.175520	-0.412995	0.360445
INDUS	-0.708027	0.595129	0.720760	0.383248	-0.356977	0.603800	-0.483725
CHAS	-0.099176	-0.007368	-0.035587	-0.121515	0.048788	-0.053929	0.175260
NOX	-0.769230	0.611441	0.668023	0.188933	-0.380051	0.590879	-0.427321
RM	0.205246	-0.209847	-0.292048	-0.355501	0.128069	-0.613808	0.695360
AGE	-0.747881	0.456022	0.506456	0.261515	-0.273534	0.602339	-0.376955
DIS	1.000000	-0.494588	-0.534432	-0.232471	0.291512	-0.496996	0.249929
RAD	-0.494588	1.000000	0.910228	0.464741	-0.444413	0.488676	-0.381626
TAX	-0.534432	0.910228	1.000000	0.460853	-0.441808	0.543993	-0.468536
PTRATIO	-0.232471	0.464741	0.460853	1.000000	-0.177383	0.374044	-0.507787
B	0.291512	-0.444413	-0.441808	-0.177383	1.000000	0.266087	0.222461
LSTAT	-0.496996	0.488676	0.543993	0.374044	0.266087	1.000000	-0.259165
MEDV	0.249929	-0.381626	-0.468536	-0.507787	0.222461	-0.259165	1.000000

Query 8: To create a heatmap of the correlation matrix:

Clause:

```
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
plt.figure(figsize=(14,7))
sns.heatmap(data.corr().abs(),annot=True)
```

Result:

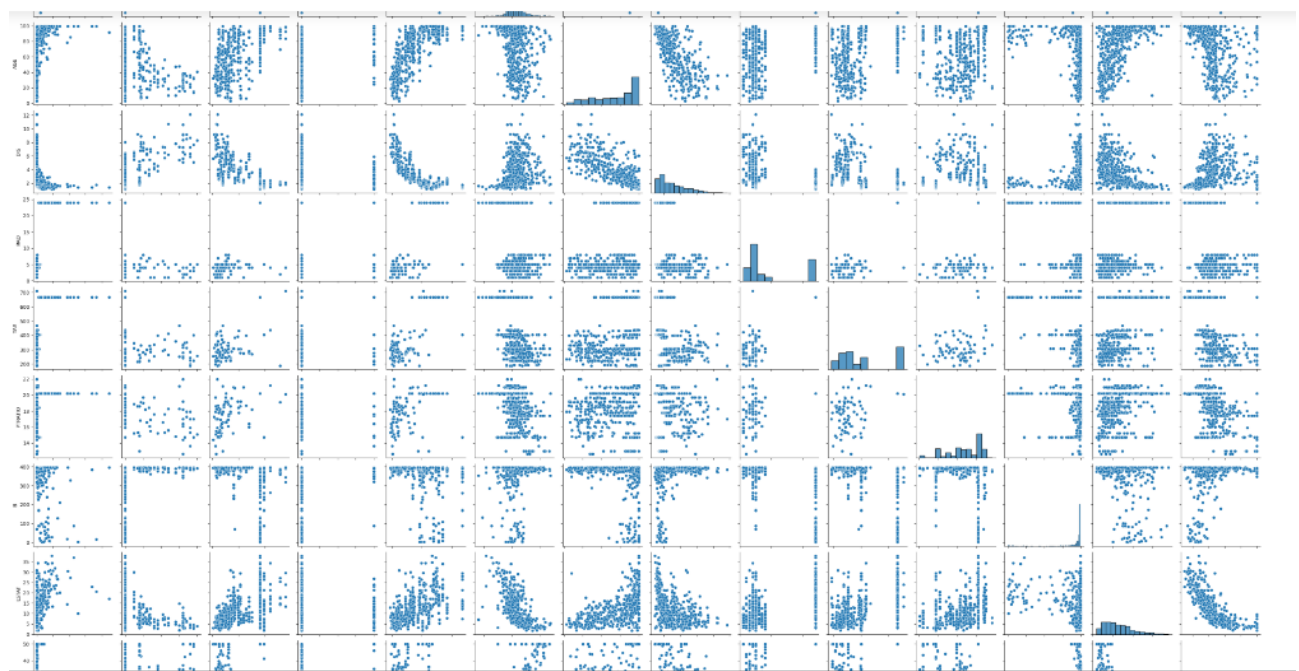


Query 9:

Clause:

```
sns.pairplot(data)
```

Result:



Query 10:

Clause:

```
x = data.drop('MEDV',axis=1)
x
```

Result:

In [17]: x

Out[17]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	5.33
...
501	0.06263	0.0	11.93	0	0.573	6.593	69.1	2.4786	1	273.0	21.0	391.99	9.67
502	0.04527	0.0	11.93	0	0.573	6.120	76.7	2.2875	1	273.0	21.0	396.90	9.08
503	0.06076	0.0	11.93	0	0.573	6.976	91.0	2.1675	1	273.0	21.0	396.90	5.64
504	0.10959	0.0	11.93	0	0.573	6.794	89.3	2.3889	1	273.0	21.0	393.45	6.48
505	0.04741	0.0	11.93	0	0.573	6.030	80.8	2.5050	1	273.0	21.0	396.90	7.88

506 rows × 13 columns

Query 11:

Clause:

```
y=data['MEDV']
y
```

Result:


```
In [18]: y
```

```
Out[18]: 0      24.0
          1      21.6
          2      34.7
          3      33.4
          4      36.2
          ...
        501     22.4
        502     20.6
        503     23.9
        504     22.0
        505     11.9
          Name: MEDV, Length: 506, dtype: float64
```

Query 12:

Clause:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test =
train_test_split(X,y,train_size=.70,test_size=0.3,random_state=50)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)
```

Result:

```
In [35]: from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(X,y,train_size=.70,test_size=0.3,random_state=50)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(354, 13)
(152, 13)
(354,)
(152,)
```

Query 13:

Clause:

```
from sklearn.linear_model import LinearRegression
m = LinearRegression()
m.fit(X_train,y_train)
```

Result:

```
In [38]: m.fit(X_train,y_train)
```

```
Out[38]: LinearRegression
LinearRegression()
```

Query 14:


Clause:

```
y_predict = m.predict(X_test)
y_pre_train= m.predict(X_train)
plt.scatter(y_train,y_pre_train,c='green',lw=1)
plt.plot([y_train.min(),y_train.max()], [y_train.min(),y_train.max()], 'k--',c='blue',lw=3)
plt.xlabel('Actual')
plt.ylabel('Predicted')
```

Result:

```
plt.plot([y_train.min(),y_train.max()], [y_train.min(),y_train.max()], 'k--',c='b]
```

```
Out[54]: Text(0, 0.5, 'Predicted')
```



Query 15:

Clause:

```
print(m.intercept_)
```

Result:

```
In [40]: print(m.intercept_)
25.469528442762886
```

Query 16:

To print coefficient

Clause:

```
print(m.coef_)
```

Result:

```
In [41]: print(m.coef_)
[-1.02245423e-01  3.11044112e-02  1.43693788e-02  1.65964577e+00
 -1.58693910e+01  5.03222436e+00 -7.16365483e-04 -1.31353766e+00
 2.55945159e-01 -1.23369178e-02 -9.03145563e-01  9.65756208e-03
 -4.22793524e-01]
```

Query 17:

To create a Pandas DataFrame that shows the coefficients of the machine learning model

Clause:

```
coefficient = pd.DataFrame(m.coef_.x.columns,columns=["Coefficient"])
coefficient
```

Result:

```
In [43]: coefficient
```

```
Out[43]:
```

Coefficient	
CRIM	-0.102245
ZN	0.031104
INDUS	0.014369
CHAS	1.659646
NOX	-15.869391
RM	5.032224
AGE	-0.000716
DIS	-1.313538
RAD	0.255945
TAX	-0.012337
PTRATIO	-0.903146
B	0.009658
LSTAT	-0.422794

Query 18:

To print the top 5 rows

Clause:

```
data.head(5)
```

Result:

```
In [44]: data.head(5)
```

```
Out[44]:
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT	MEDV
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	4.98	24.0
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	9.14	21.6
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	4.03	34.7
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	2.94	33.4
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	5.33	36.2

Query 19:

To predict the median house price for the house

Clause:

```
m.predict([[0.02729,0.0,7.07,0,0.469,7.185,61.1,4.9671,2,242.0,17.8,392.83,4.03]])
```

Result:

```
In [45]: m.predict([[0.02729,0.0,7.07,0,0.469,7.185,61.1,4.9671,2,242.0,17.8,392.83,4.03]])
C:\ProgramData\anaconda3\lib\site-packages\sklearn\base.py:420: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
Out[45]: array([31.25416184])
```

Query 20:

To predict the median house price for the house

Clause:

```
m.predict([[0.06905,0.0,2.18,0,0.458,7.147,54.2,6.0622,3,222.0,18.7,396.90,5.33]])
```

Result:

```
In [46]: m.predict([[0.06905,0.0,2.18,0,0.458,7.147,54.2,6.0622,3,222.0,18.7,396.90,5.33]])
C:\ProgramData\anaconda3\lib\site-packages\sklearn\base.py:420: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names
  warnings.warn(
Out[46]: array([28.90897962])
```

Query 21:

Evaluate the performance of the model

Clause:

```
from sklearn import metrics
print('MAE :',metrics.mean_absolute_error(y_test,y_predict))
print('MSE :',metrics.mean_squared_error(y_test,y_predict))
```

Result:

```
In [48]: print('MAE :',metrics.mean_absolute_error(y_test,y_predict))
print('MSE :',metrics.mean_squared_error(y_test,y_predict))

MAE : 3.678977534499423
MSE : 33.868033996670015
```