

**Institute of Information Technology (IIT)**  
**Jahangirnagar University**



**Lab Report: 06**

Submitted by:

Name: Zannat Hossain Tamim

Roll No:1970

Lab Date:22.08.23

Submission Date: 28.08.23

## Lab Report # Day 06

### Query 1:

Import libraries , read CSV file, and print the file

Clause:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
df= pd.read_csv('diabetes.csv')
df
```

Result :

In [126]: df

Out[126]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...	...	...	...	...	...	...	...	...	...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

Query 2: Print the top 4 rows.

Clause:

```
df.head(4)
```

Result :

In [127]: df.head(4)

Out[127]:

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0

**Query 3:** Print the last 4 rows of the data frame.

**Clause:**

```
df.tail()
```

**Result :**

```
In [128]: df.tail(4)
```

```
Out[128]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

**Query 4:**

To get the shape of the DataFrame

**Clause:**

```
df.shape
```

**Result :**

```
In [183]: df.shape
```

```
Out[183]: (768, 9)
```

**Query 5:**

To get the information of the DataFrame

**Clause:**

```
df.info()
```

**Result :**

```
In [129]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Pregnancies           768 non-null   int64  
 1   Glucose               768 non-null   int64  
 2   BloodPressure         768 non-null   int64  
 3   SkinThickness         768 non-null   int64  
 4   Insulin               768 non-null   int64  
 5   BMI                   768 non-null   float64 
 6   DiabetesPedigreeFunction 768 non-null   float64 
 7   Age                  768 non-null   int64  
 8   Outcome               768 non-null   int64  
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```

## Query 6:

To calculate and print a summary of statistical data for each column in the DataFrame

**Clause:**

```
df.describe()
```

**Result :**

```
Out[184]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

## Query 7:

To count the number of missing values in each column of a DataFrame

**Clause:**

```
df.isnull().sum()
```

**Result:**

```
In [186]: df.isnull().sum()
```

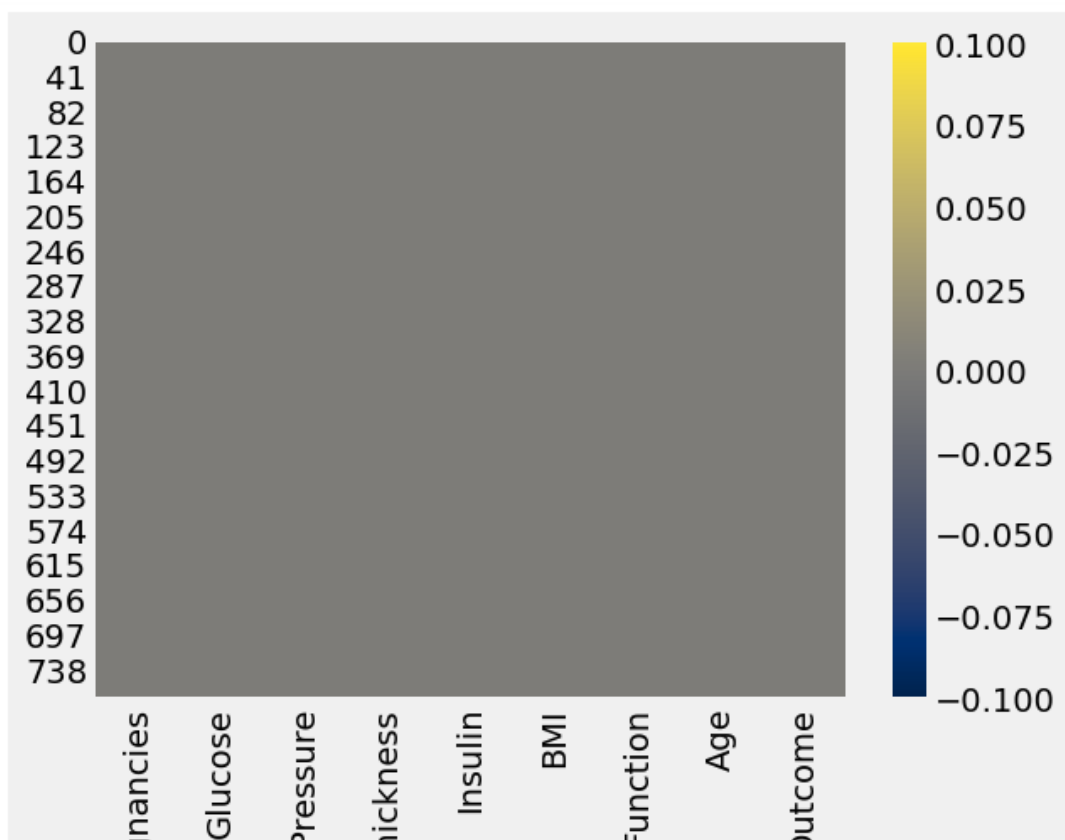
```
Out[186]: Pregnancies      0
          Glucose          0
          BloodPressure    0
          SkinThickness    0
          Insulin          0
          BMI              0
          DiabetesPedigreeFunction  0
          Age              0
          Outcome          0
          dtype: int64
```

**Query 8 :** Create a heatmap of the missing values in the data frame.

**Clause:**

```
sns.heatmap(df.isnull(),cmap='cividis')
```

**Result:**




**Query 9 :**

**Clause:**

```
p=df.Outcome.value_counts()
p.plot(kind="bar")
```

**Result:**

ut[206]: <Axes: >



**Query 10:** Calculate the correlation coefficient between each pair of columns in the data frame and store it in the variable `cor`

**Clause:**

```
cor=df.corr()  
print(cor)
```

**Result:**

	Pregnancies	Glucose	BloodPressure	SkinThickness	\
Pregnancies	1.000000	0.129459	0.141282	-0.081672	
Glucose	0.129459	1.000000	0.152590	0.057328	
BloodPressure	0.141282	0.152590	1.000000	0.207371	
SkinThickness	-0.081672	0.057328	0.207371	1.000000	
Insulin	-0.073535	0.331357	0.088933	0.436783	
BMI	0.017683	0.221071	0.281805	0.392573	
DiabetesPedigreeFunction	-0.033523	0.137337	0.041265	0.183928	
Age	0.544341	0.263514	0.239528	-0.113970	
Outcome	0.221898	0.466581	0.065068	0.074752	

	Insulin	BMI	DiabetesPedigreeFunction	\
Pregnancies	-0.073535	0.017683	-0.033523	
Glucose	0.331357	0.221071	0.137337	
BloodPressure	0.088933	0.281805	0.041265	
SkinThickness	0.436783	0.392573	0.183928	
Insulin	1.000000	0.197859	0.185071	
BMI	0.197859	1.000000	0.140647	
DiabetesPedigreeFunction	0.185071	0.140647	1.000000	
Age	-0.042163	0.036242	0.033561	
Outcome	0.130548	0.292695	0.173844	


	Age	Outcome
Pregnancies	0.544341	0.221898
Glucose	0.263514	0.466581
BloodPressure	0.239528	0.065068
SkinThickness	-0.113970	0.074752
Insulin	-0.042163	0.130548
BMI	0.036242	0.292695
DiabetesPedigreeFunction	0.033561	0.173844

**Query 11:** To create a heatmap of the correlation matrix:

**Clause:**

```
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
plt.figure(figsize=(12,6))
sns.heatmap(df.corr().abs(),annot=True)
```

**Result:**




## Query 12:

Plot pairwise relationships between variables of the dataset

**Clause:**

```
sns.pairplot(df)
```

**Result:**



### Query 13: Number of *Outcome*

Clause:

```
df.groupby('Outcome').size()
```

Result:

```
In [166]: df.groupby('Outcome').size()
Out[166]: Outcome
0      500
1      268
dtype: int64
```

### Query 14:

Standardize the variables

Clause:

```
from sklearn.preprocessing import StandardScaler
s=StandardScaler()
s.fit(df.drop('Outcome',axis=1))
```

Result:



```
In [132]: s.fit(df.drop('Outcome',axis=1))
Out[132]: StandardScaler
StandardScaler()
```

## Query 15:

Standardize the variables

**Clause:**

```
sf= s.transform(df.drop('Outcome',axis=1))
df2= pd.DataFrame(sf,columns=df.columns[:-1])
df2
```

**Result:**

```
In [135]: df2
```

```
Out[135]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	0.639947	0.848324	0.149641	0.907270	-0.692891	0.204013	0.468492	1.425995
1	-0.844885	-1.123396	-0.160546	0.530902	-0.692891	-0.684422	-0.365061	-0.190672
2	1.233880	1.943724	-0.263941	-1.288212	-0.692891	-1.103255	0.604397	-0.105584
3	-0.844885	-0.998208	-0.160546	0.154533	0.123302	-0.494043	-0.920763	-1.041549
4	-1.141852	0.504055	-1.504687	0.907270	0.765836	1.409746	5.484909	-0.020496
...	...	...	...	...	...	...	...	...
763	1.827813	-0.622642	0.356432	1.722735	0.870031	0.115169	-0.908682	2.532136
764	-0.547919	0.034598	0.046245	0.405445	-0.692891	0.610154	-0.398282	-0.531023
765	0.342981	0.003301	0.149641	0.154533	0.279594	-0.735190	-0.685193	-0.275760
766	-0.844885	0.159787	-0.470732	-1.288212	-0.692891	-0.240205	-0.371101	1.170732
767	-0.844885	-0.873019	0.046245	0.656358	-0.692891	-0.202129	-0.473785	-0.871374

768 rows × 8 columns

## Query 16:

**Clause:**

```
df2.head(4))
```

**Result:**

```
[136]: df2.head(4)
```

```
t[136]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	0.639947	0.848324	0.149641	0.907270	-0.692891	0.204013	0.468492	1.425995
1	-0.844885	-1.123396	-0.160546	0.530902	-0.692891	-0.684422	-0.365061	-0.190672
2	1.233880	1.943724	-0.263941	-1.288212	-0.692891	-1.103255	0.604397	-0.105584
3	-0.844885	-0.998208	-0.160546	0.154533	0.123302	-0.494043	-0.920763	-1.041549

## Query 17:

Train test split and using KNN

**Clause:**

```

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df.loc[:, df.columns != 'Outcome'],
                                                    df['Outcome'], stratify=df['Outcome'], random_state=101)
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=1)
knn

```

**Result:**

```

In [140]: knn
Out[140]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=1)

```

## Query 18:

**Clause:**

```
knn.fit(X_train, y_train)
```

**Result:**

```

In [141]: knn.fit(X_train, y_train)
Out[141]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=1)

```

## Query 19:

**Clause:**

```

pred=knn.predict(X_test)
pred

```

**Result:**

```

In [143]: pred
Out[143]: array([1, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1,
0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 1,
0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1,
0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0,
1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0,
1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0,
0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)

```

## Query 20:

Prediction and evaluation

**Clause:**

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
confusion_matrix(y_test ,pred)
```

**Result:**

```
In [145]: confusion_matrix(y_test ,pred)
```

```
Out[145]: array([[98, 27],
                 [30, 37]], dtype=int64)
```

### Query 21:

Prediction and evaluation

**Clause:**

```
print(classification_report(y_test ,pred))
```

**Result:**

```
In [146]: print(classification_report(y_test ,pred))
```

	precision	recall	f1-score	support
0	0.77	0.78	0.77	125
1	0.58	0.55	0.56	67
accuracy			0.70	192
macro avg	0.67	0.67	0.67	192
weighted avg	0.70	0.70	0.70	192

### Query 22:

Find the error rate to choose a K value

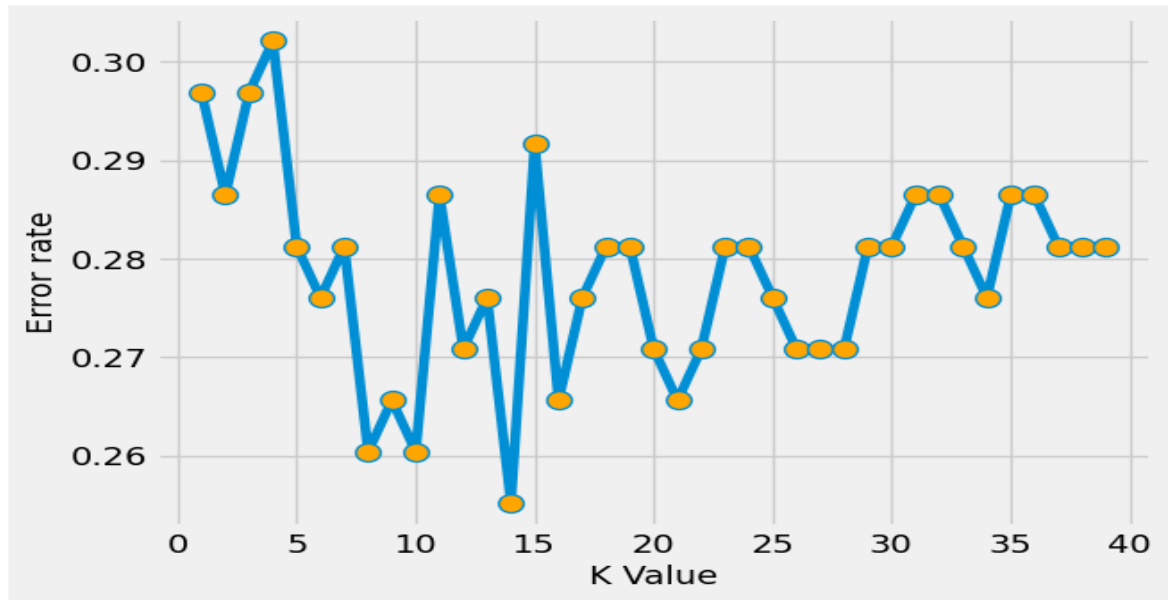
**Clause:**

```
err=[]
for i in range(1,40):
    knn= KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train, y_train)
    pred_i=knn.predict(X_test)
    err.append(np.mean( pred_i!= y_test))

plt.plot(range(1,40),err,marker='o',markerfacecolor='orange',markersize=10)
plt.xlabel('K Value')
plt.ylabel('Error rate')
```

**Result:**

Out[148]: Text(0, 0.5, 'Error rate')



### Query 23:

For K=1,

**Clause:**

```
knn = KNeighborsClassifier(n_neighbors=1)
knn.fit(X_train, y_train)
pred2=knn.predict(X_test)
print('For K=1, \n')
print('confusion matrix is : ')
print(confusion_matrix(y_test ,pred),'\n')
print('classification report is : ')
print(classification_report(y_test ,pred))
```

**Result:**

For K=1,

confusion matrix is :

```
[[98 27]
 [30 37]]
```

classification report is :

	precision	recall	f1-score	support
0	0.77	0.78	0.77	125
1	0.58	0.55	0.56	67
accuracy			0.70	192
macro avg	0.67	0.67	0.67	192
weighted avg	0.70	0.70	0.70	192


### Query 24:

Confusion Matrix Display for K=1

Clause:

```
from sklearn.metrics import ConfusionMatrixDisplay
c=confusion_matrix(y_test ,pred2)
c2=ConfusionMatrixDisplay(c)
c2.plot()
plt.grid(False)
```

Result:



### Query 25:

For K=14,

Clause:

```
knn = KNeighborsClassifier(n_neighbors=14)
knn.fit(X_train, y_train)
pred3=knn.predict(X_test)
```

```
print('For K=14, \n')
print('confusion matrix is : ')
print(confusion_matrix(y_test ,pred3),'\n')
print('classification report is : ')
print(classification_report(y_test ,pred3))
```

### Result:

```
For K=14,

confusion matrix is :
[[117  8]
 [ 41 26]]

classification report is :
              precision    recall  f1-score   support

     0           0.74       0.94       0.83        125
     1           0.76       0.39       0.51         67

   accuracy                   0.74        192
  macro avg           0.75       0.66       0.67        192
 weighted avg           0.75       0.74       0.72        192
```


### Query 26:

Confusion Matrix Display for K=14

### Clause:

```
from sklearn.metrics import ConfusionMatrixDisplay
c=confusion_matrix(y_test ,pred3)
c2=ConfusionMatrixDisplay(c)
c2.plot()
plt.grid(False)
```

### Result:



### Query 27:

For k=10

#### Clause:

```
knn = KNeighborsClassifier(n_neighbors=10)
knn.fit(X_train, y_train)
pred5=knn.predict(X_test)
print('For K=8, \n')
print('confusion matrix is : ')
print(confusion_matrix(y_test ,pred5),'\n')
print('classification report is : ')
print(classification_report(y_test ,pred5))
```

#### Result:

For K=10,

confusion matrix is :

```
[[114  11]
 [ 39  28]]
```

classification report is :

	precision	recall	f1-score	support
0	0.75	0.91	0.82	125
1	0.72	0.42	0.53	67
accuracy			0.74	192
macro avg	0.73	0.66	0.67	192
weighted avg	0.74	0.74	0.72	192

### Query 28:


Confusion Matrix Display for K=14

**Clause:**

```
from sklearn.metrics import ConfusionMatrixDisplay
c=confusion_matrix(y_test,pred5)
c2=ConfusionMatrixDisplay(c)
c2.plot()
plt.grid(False)
```

**Result:**





### Query 29:

Check the accuracy

Clause:

```
print('Accuracy of K-NN classifier on training set: {:.2f}'.format(knn.score(X_train, y_train)))
print('Accuracy of K-NN classifier on test set: {:.2f}'.format(knn.score(X_test, y_test)))
```

Result:

```
In [181]: print('Accuracy of K-NN classifier on training set: {:.2f}'.format(knn.score(X_train, y_train)))
          print('Accuracy of K-NN classifier on test set: {:.2f}'.format(knn.score(X_test, y_test)))

Accuracy of K-NN classifier on training set: 0.79
Accuracy of K-NN classifier on test set: 0.74
```

### Query 30:

For K=9

Clause:

```
knn = KNeighborsClassifier(n_neighbors=9)
knn.fit(X_train, y_train)
pred6=knn.predict(X_test)
print('For K=9, \n')
print('confusion matrix is : ')
print(confusion_matrix(y_test, pred6), '\n')
print('classification report is : ')

```

```
print(classification_report(y_test ,pred6))
```

### Result:

For K=9,

confusion matrix is :

```
[[112  13]
 [ 22  45]]
```

classification report is :

	precision	recall	f1-score	support
0	0.84	0.90	0.86	125
1	0.78	0.67	0.72	67
accuracy			0.82	192
macro avg	0.81	0.78	0.79	192
weighted avg	0.81	0.82	0.81	192


### Query 31:

Confusion Matrix Display for K=9

#### Clause:

```
from sklearn.metrics import ConfusionMatrixDisplay
c=confusion_matrix(y_test ,pred6)
c2=ConfusionMatrixDisplay(c)
c2.plot()
plt.grid(False)
```

### Result:



### Query 32:

Check the accuracy

Clause:

```
print('Accuracy of K-NN classifier on training set: {:.2f}'.format(knn.score(X_train, y_train)))  
print('Accuracy of K-NN classifier on test set: {:.2f}'.format(knn.score(X_test, y_test)))
```

Result:

```
In [193]: print('Accuracy of K-NN classifier on training set: {:.2f}'.format(knn.score(X_train, y_train)))  
          print('Accuracy of K-NN classifier on test set: {:.2f}'.format(knn.score(X_test, y_test)))  
  
          Accuracy of K-NN classifier on training set: 0.77  
          Accuracy of K-NN classifier on test set: 0.82
```

### Query 33:

Check the accuracy

Clause:

```
from sklearn.metrics import roc_auc_score  
roc_auc_score(y_test, pred6)
```

Result:

```
In [194]: from sklearn.metrics import roc_auc_score  
          roc_auc_score(y_test, pred6)  
  
Out[194]: 0.783820895522388
```